

MOSAIC'23 PS-2

IMAGE- INPAINTING MODEL

@442|PURAV HARAN|PRASOON SAHAY|PRIYANSHU

OVERVIEW

- > Setup, installations and imports
- > Preparing dataset
- > Partial convolution based auto-encoder-decoder model
- > Training and testing

SETUP,
INSTALLATIONS
&
IMPORTS

Setups, Installations and Imports

```
[ ] !pip install tensorflow-gpu==2.0

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
ERROR: Could not find a version that satisfies the requirement tensorflow-gpu==2.0 (from versions: 2.5.0, 2.5.1, 2.5.2, 2.5.3, 2.6.0, 2.6.1, 2.6.2, 2.6.3, 2.6.4, 2.6.5, 2.7.0rc0, 2.7.0rc1, 2.7.0, 2.7.1, 2.7.2, 2
ERROR: No matching distribution found for tensorflow-gpu==2.0
```

```
[ ] !pip install wandb -q
```

```
▶ import tensorflow as tf
print(tf.__version__)
from tensorflow import keras

print('[INFO]', tf.config.experimental.list_physical_devices('GPU')[0])
```

```
⇨ 2.12.0
[INFO] PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')
```

```
[ ] import wandb
from wandb.keras import WandbCallback

wandb.login()
```

```
wandb: Currently logged in as: purav-haran1305. Use `wandb login --relogin` to force relogin
True
```

```
[ ] import os
import cv2
import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline

from mpl_toolkits.axes_grid1 import ImageGrid
```

PREPARING DATASET

We've taken a CIFAR-10 dataset to train the model

▼ Visualization of Cifar 10

```
▶ ## Get first 32 images as samples
sample_images = x_train[:32]
sample_labels = y_train[:32]

fig = plt.figure(figsize=(16., 8.))
grid = ImageGrid(fig, 111,
                  nrows_ncols=(4, 8), # creates 2x2 grid of axes
                  axes_pad=0.3, # pad between axes in inch.
                  )

for ax, image, label in zip(grid, sample_images, sample_labels):
    ax.imshow(image)
    ax.set_title(label[0])
```

We've generated the data with Patch Augmentation

```
## Ref: https://stanford.edu/~shervine/blog/keras-how-to-generate-data-on-the-fly.
class createAugment(keras.utils.Sequence):
    'Generates data for Keras'
    def __init__(self, X, y, batch_size=32, dim=(32, 32), n_channels=3, shuffle=True):
        'Initialization'
        self.batch_size = batch_size
        self.X = X
        self.y = y
        self.dim = dim
        self.n_channels = n_channels
        self.shuffle = shuffle

        self.on_epoch_end()

    def __len__(self):
        'Denotes the number of batches per epoch'
        return int(np.floor(len(self.X) / self.batch_size))

    def __getitem__(self, index):
        'Generate one batch of data'
        # Generate indexes of the batch
        indexes = self.indexes[index * self.batch_size:(index+1)*self.batch_size]

        # Generate data
        X_inputs, y_output = self.__data_generation(indexes)
        return X_inputs, y_output

    def on_epoch_end(self):
        'Updates indexes after each epoch'
        self.indexes = np.arange(len(self.X))
        if self.shuffle:
            np.random.shuffle(self.indexes)

    def __data_generation(self, idxs):
        # Masked_images is a matrix of masked images used as input
        Masked_images = np.empty((self.batch_size, self.dim[0], self.dim[1], self.n_channels)) # Masked image
        # Mask_batch is a matrix of binary masks used as input
        Mask_batch = np.empty((self.batch_size, self.dim[0], self.dim[1], self.n_channels)) # Binary Masks
        # y_batch is a matrix of original images used for computing error from reconstructed image
        y_batch = np.empty((self.batch_size, self.dim[0], self.dim[1], self.n_channels)) # Original image

        ## Iterate through random indexes
        for i, idx in enumerate(idxs):
            image_copy = self.X[idx].copy()

            ## Get mask associated to that image
            masked_image, mask = self.__createMask(image_copy)

            Masked_images[i,] = masked_image/255
            Mask_batch[i,] = mask/255
            y_batch[i,] = self.y[idx]/255

        ## Return mask as well because partial convolution require the same.
        return [Masked_images, Mask_batch], y_batch

    def __createMask(self, img):
        ## Prepare masking matrix
        mask = np.full((32,32,3), 255, np.uint8) ## White background
        for _ in range(np.random.randint(1, 10)):
            # Get random x locations to start line
            x1, x2 = np.random.randint(1, 32), np.random.randint(1, 32)
            # Get random y locations to start line
            y1, y2 = np.random.randint(1, 32), np.random.randint(1, 32)
```

```

    masked_image[mask==0] = 255
    return masked_image, mask

[11] ## Prepare training and testing mask-image pair generator
traingen = createAugment(x_train, x_train)
testgen = createAugment(x_test, x_test, shuffle=False)

[12] # Legend: Original Image | Mask generated | Masked Image

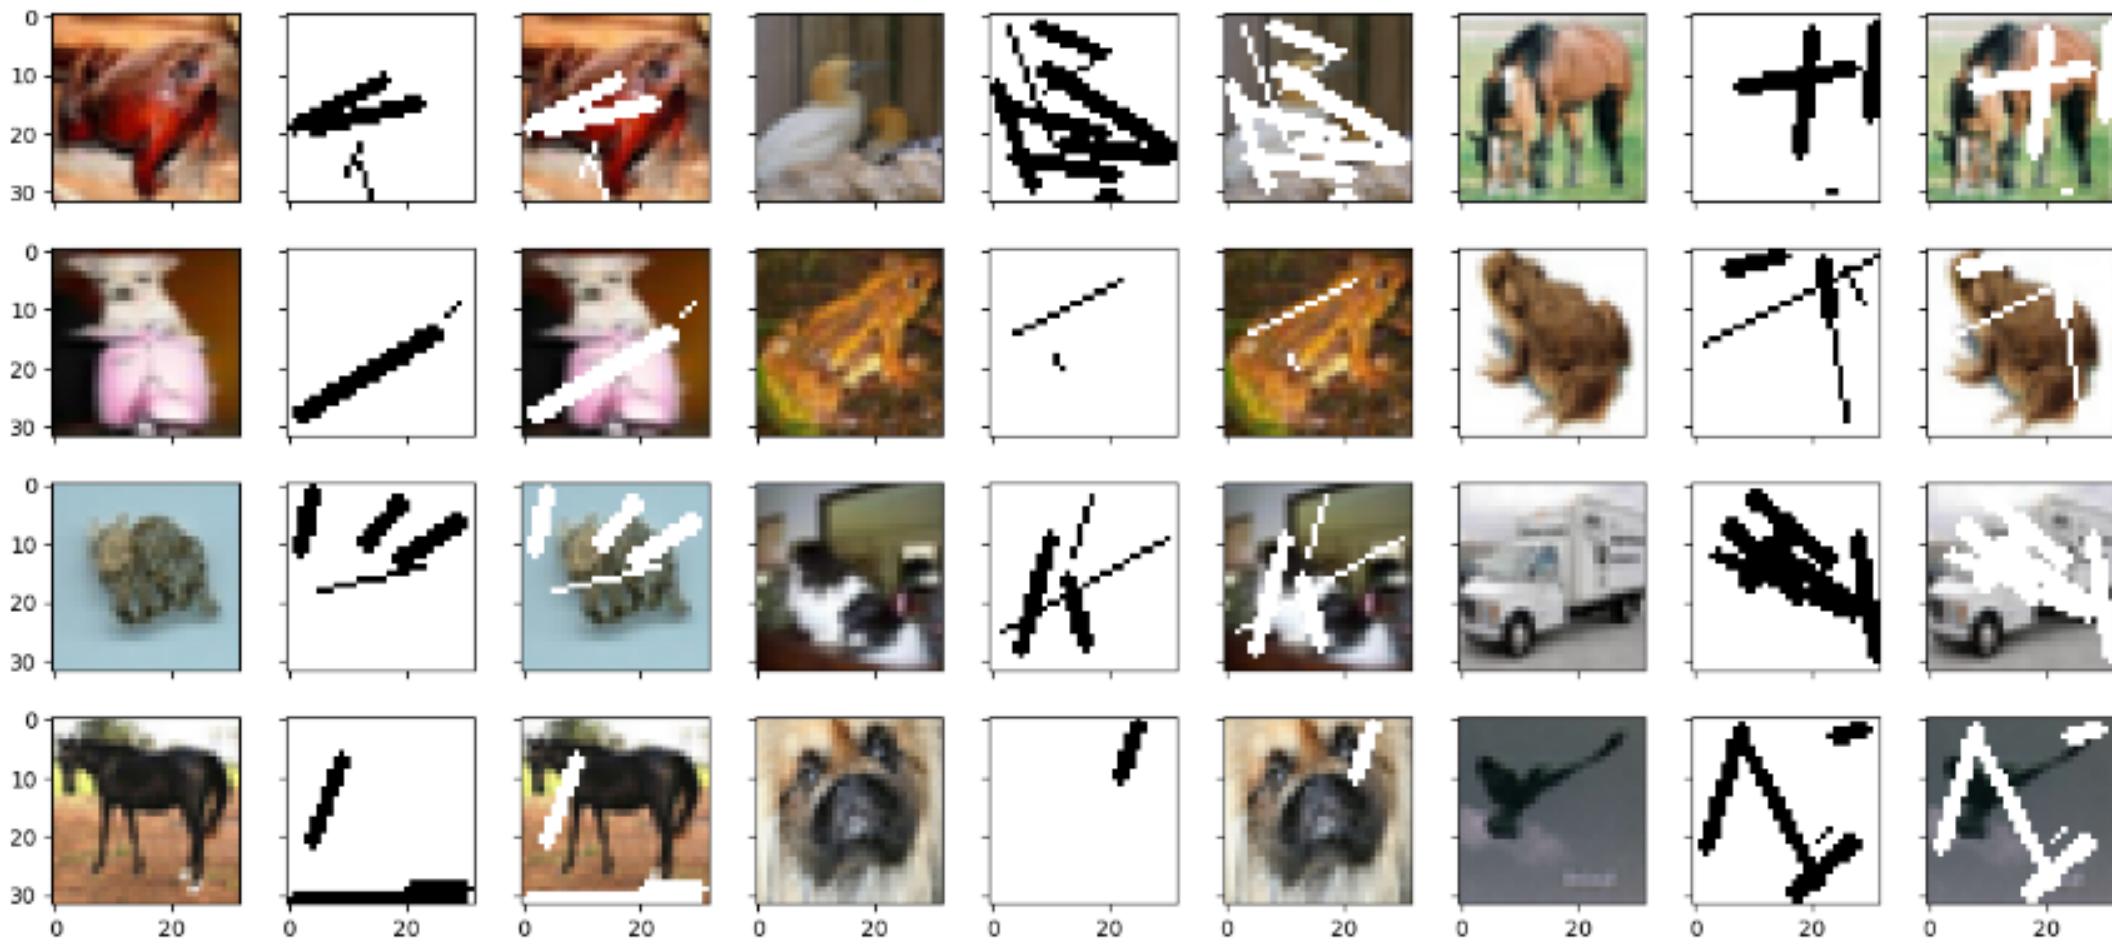
## Examples
sample_idx = 90 ## Change this to see different batches

[masked_images, masks], sample_labels = traingen[sample_idx]
sample_images = [None]*(len(masked_images)+len(masks)+len(sample_labels))
sample_images[::3] = sample_labels
## masks[]
sample_images[1::3] = masks
sample_images[2::3] = masked_images

fig = plt.figure(figsize=(17., 8.))
grid = ImageGrid(fig, 111, ## similar to subplot(111)
                 nrows_ncols=(4, 9), ## creates 2x2 grid of axes
                 axes_pad=0.3, ## pad between axes in inch.
                 )
for ax, image in zip(grid, sample_images):
    ax.imshow(image)

plt.show()

```



Partial convolution based auto-encoder-decoder model

```
✓ [13] #! utils is present in the cloned repo. Visit repo for the implementation of PConv2D.
from utils.pconv_layer import PConv2D

✓ [14] #! For more information into formulation: https://www.youtube.com/watch?v=AZr640xshLo
#! Metric
def dice_coef(y_true, y_pred):
    y_true_f = keras.backend.flatten(y_true)
    y_pred_f = keras.backend.flatten(y_pred)
    intersection = keras.backend.sum(y_true_f * y_pred_f)
    return (2. * intersection) / (keras.backend.sum(y_true_f + y_pred_f))

✓ [15] class InpaintingModel:
    ...
    Build UNET like model for image inpaining task.
    ...

    def prepare_model(self, input_size=(32,32,3)):
        input_image = keras.layers.Input(input_size)
        input_mask = keras.layers.Input(input_size, name='encoder_input')

        conv1, mask1, conv2, mask2 = self._encoder_layer(32, input_image, input_mask, ['conv1', 'conv2'])
        conv3, mask3, conv4, mask4 = self._encoder_layer(64, conv2, mask2, ['conv3', 'conv4'])
        conv5, mask5, conv6, mask6 = self._encoder_layer(128, conv4, mask4, ['conv5', 'conv6'])
        conv7, mask7, conv8, mask8 = self._encoder_layer(256, conv6, mask6, ['conv7', 'encoder_output'])

        conv9, mask9, conv10, mask10 = self._decoder_layer(256, 128, conv8, mask8, conv7, mask7, ['conv9', 'conv10'])
        conv11, mask11, conv12, mask12 = self._decoder_layer(128, 64, conv10, mask10, conv5, mask5, ['conv11', 'conv12'])
        conv13, mask13, conv14, mask14 = self._decoder_layer(64, 32, conv12, mask12, conv3, mask3, ['conv13', 'conv14'])
        conv15, mask15, conv16, mask16 = self._decoder_layer(32, 3, conv14, mask14, conv1, mask1, ['conv15', 'decoder_output'])

        outputs = keras.layers.Conv2D(3, (3, 3), activation='sigmoid', padding='same')(conv16)

        return keras.models.Model(inputs=[input_image, input_mask], outputs=[outputs])

    def _encoder_layer(self, filters, in_layer, in_mask, names):
        conv1, mask1 = PConv2D(32, (3,3), strides=1, padding='same', name=names[0])([in_layer, in_mask])
        conv1 = keras.activations.relu(conv1)

        conv2, mask2 = PConv2D(32, (3,3), strides=2, padding='same', name=names[1])([conv1, mask1])
        # conv2 = keras.layers.BatchNormalization()(conv2, training=True)
        conv2 = keras.activations.relu(conv2)

        return conv1, mask1, conv2, mask2

    def _decoder_layer(self, filter1, filter2, in_img, in_mask, share_img, share_mask, names):
        up_img = keras.layers.UpSampling2D(size=(2,2))(in_img)
        up_mask = keras.layers.UpSampling2D(size=(2,2))(in_mask)
        concat_img = keras.layers.concatenate([share_img, up_img])
        concat_mask = keras.layers.concatenate([share_mask, up_mask])

        conv1, mask1 = PConv2D(filter1, (3,3), padding='same', name=names[0])([concat_img, concat_mask])
        conv1 = keras.activations.relu(conv1)

        conv2, mask2 = PConv2D(filter2, (3,3), padding='same', name=names[1])([conv1, mask1])
        # conv2 = keras.layers.BatchNormalization()(conv2)
        conv2 = keras.activations.relu(conv2)

        return conv1, mask1, conv2, mask2
```

Partial Convolution performs the normalization of the output to adjust for the fraction of missing data. A partial convolution layer comprises masked and re-normalized convolution operation followed by a mask-update setup.

TRAINING

Train

```
[ ] wandb.init(  
    # set the wandb project where this run will be logged  
    project="my-awesome-project"  
  
)  
  
Tracking run with wandb version 0.14.0  
Run data is saved locally in /content/deepimageinpainting/wandb/run-20230404_220934-7xto16n2  
Syncing run different-pine-6 to Weights & Biases \(docs\)  
View project at https://wandb.ai/purav-haran1305/my-awesome-project  
View run at https://wandb.ai/purav-haran1305/my-awesome-project/runs/7xto16n2  
Display W&B run
```

```
▶ class PredictionLogger(tf.keras.callbacks.Callback):  
    def __init__(self):  
        super(PredictionLogger, self).__init__()  
  
    def on_epoch_end(self, logs, epoch):  
        sample_idx = 54  
        [masked_images, masks], sample_labels = testgen[sample_idx]  
  
        m_images = []  
        binary_masks = []  
        predictions = []  
        labels = []  
  
        for i in range(32):  
            inputs = [masked_images[i].reshape((1,)+masked_images[i].shape), masks[i].reshape((1,)+masks[i].shape)]  
            impainted_image = model.predict(inputs)  
  
            m_images.append(masked_images[i])  
            binary_masks.append(masks[i])  
            predictions.append(impainted_image.reshape(impainted_image.shape[1:])))  
            labels.append(sample_labels[i])  
  
        wandb.log({"masked_images": [wandb.Image(m_image)  
                                    for m_image in m_images]})  
        wandb.log({"masks": [wandb.Image(mask)  
                            for mask in binary_masks]})  
        wandb.log({"predictions": [wandb.Image(inpainted_image)  
                                 for inpainted_image in predictions]})  
        wandb.log({"labels": [wandb.Image(label)  
                           for label in labels]})
```

TESTING

Testing on images

```
## Legend: Original Image | Mask generated | Inpainted Image | Ground Truth  
  
## Examples  
rows = 32  
sample_idx = 54  
[masked_images, masks], sample_labels = testgen[sample_idx]  
  
fig, axes = plt.subplots(nrows=rows, ncols=4, figsize=(8, 2*rows))  
  
for i in range(32):  
    inputs = [masked_images[i].reshape((1,)+masked_images[i].shape), masks[i].reshape((1,)+masks[i].shape)]  
    impainted_image = model.predict(inputs)  
    axes[i][0].imshow(masked_images[i])  
    axes[i][1].imshow(masks[i])  
    axes[i][2].imshow(impainted_image.reshape(impainted_image.shape[1:])))  
    axes[i][3].imshow(sample_labels[i])  
  
plt.show()
```



```
1/1 [=====] - 0s 22ms/step  
1/1 [=====] - 0s 27ms/step  
1/1 [=====] - 0s 26ms/step  
1/1 [=====] - 0s 26ms/step  
1/1 [=====] - 0s 27ms/step  
1/1 [=====] - 0s 25ms/step  
1/1 [=====] - 0s 24ms/step  
1/1 [=====] - 0s 24ms/step
```

SAVING THE MODEL BY:
MODEL.SAVE('/CONTENT/SAMPLE
_DATA/MODEL1.H5')

Thanks Y'all

